

Hands on Machine Learning for Fluid Dynamics

7 – 11 February 2022



Lecture 4 **Tutorial Exercise 1:** **Regression with Uncertainties**

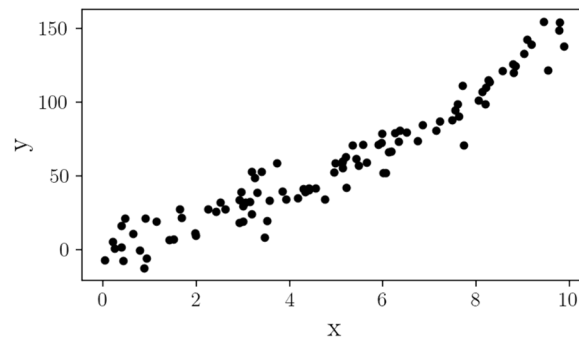
Miguel Alfonso Mendez
mendez@vki.ac.be

Slide 1

Consider Today's dataset



```
# Generate data
n_p=100
np.random.seed(10)
x_s=np.random.uniform(0,10,n_p)
y_s=2*x_s+2+np.random.normal(loc=0,scale=10,size=len(x_s))+x_s**2
```



Slide 2

Step 1: Fitting Lines and Parabolas



Build two Python functions to perform the linear regression of the data using the analytic solution we have derived today:

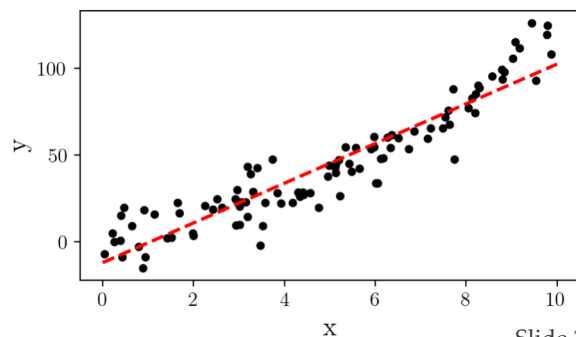
$$\mathbf{y} = \mathbf{X}\mathbf{w} \longrightarrow \mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$\mathbf{X} = \begin{bmatrix} 1 & x_0 \\ \vdots & \vdots \\ 1 & x_{n_p-1} \end{bmatrix} \in \mathbb{R}^{n_p \times 2}$ $\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$

```
def line_s_fit(x,y):  
    # take input vectors x,y and fit a line  
    # with coefficients w_0 and w_1  
  
    return w
```

Validate your results with numpy's polyfit. Then, build another function that uses a second order polynomial.

Note that the exact same approach can be used to fit higher order polynomials.



Slide 3

Step 2: Robust (Regularized) Fitting



As we will see more in details in the Lecture 9, it is often a good idea to increase the robustness of the regression by adding a regularization. Let's see it here from a numerical stability point of view.

We are essentially solving the following linear system

$$\mathbf{H}\mathbf{w} = \mathbf{b} \quad \text{with} \quad \mathbf{H} = \mathbf{X}^T\mathbf{X} \quad \mathbf{b} = \mathbf{X}^T\mathbf{y}$$

In a linear system, perturbations in the RHS leads to perturbation in the solution following the classic inequality:

$$\frac{\|\delta\mathbf{w}\|}{\|\mathbf{w}\|} \leq \underbrace{\kappa(\mathbf{H})}_{\text{Condition number}} \frac{\|\delta\mathbf{b}\|}{\|\mathbf{b}\|}$$

Condition number of H:
ratio of largest to smallest eigenvalue

$$\kappa(\mathbf{H}) = \lambda_M / \lambda_m$$

To regularize means to improve (decrease) the condition number.

A simple approach (Ridge regression) consist in adding a small number in the diagonal.

$$\mathbf{H}' = \mathbf{H} + \alpha\mathbf{I}$$

Slide 4

Step 2: Robust (Regularized) Fitting



Here are some questions for you

1. What is the condition number of the Hessian for the linear and the quadratic problem?
2. Derive a formula to link the Ridge coefficient α to the largest condition number you tolerate
3. Prepare two python functions to compute the weights for the linear and quadratic fit taking as used input a limiting condition number.

```
def Quadratic_s_fit_red(x,y,k_A_l=1000):  
    # take input vectors x,y and fit a line  
    # with coefficients w_0 and w_1  
    n_p=len(x)  
    X=np.zeros((n_p,3))  
    X[:,0]=  
    X[:,1]=  
    X[:,2]=  
    # Compute the Hessian  
    H=X.T@(X)  
    # compute the eigenvalues, the largest (l_M) and the smallest (l_m)  
    Lambd=np.linalg.eig(H); l_M=np.max(Lambd[0]); l_m=np.min(Lambd[0]);  
    alpha=  
    H_p=H+alpha*np.identity(np.shape(H)[0])  
    w=np.linalg.inv(H_p).dot(X.T).dot(y)  
    return w
```

Note: check that the regularized Hessian has the condition number you set! (see Python's function `np.linalg.cond`)

Slide 5

Step 3: Uncertainty Analysis



We are interested in analyzing the uncertainties of the regression for both regularized and un-regularized regressions. We follow an ensemble approach to compute uncertainties directly from the posterior population

- 1) Write a function that computes the regression using the analytical tools for n_e times, using each time 70% of the data for training and 30% of the data for validation.
- 2) From the n_e results, compute the average in sample and out of sample errors.
- 3) Then, given an arbitrary grid x_g , provide the mean prediction and the associated uncertainties using the ensemble statistics:

$$\mu_y(x) = \frac{1}{n_e} \sum_{j=0}^{n_e-1} y_j(x) \quad \sigma_y^2(x) = \frac{1}{n_e} \sum_{j=0}^{n_e-1} (y_j(x) - \mu_y(x))^2$$

- 4) For each regression, store the resulting weights. Then look at the distribution (histogram) of each weight.

Slide 6

```

def Boot_Strap_Model_2_esamble(X,y,x_g,n_e=500,tp=0.3):
    Pop_Y=np.zeros((len(x_g),n_e)) # posterior population
    mu_y=np.zeros(len(x_g)) # mean prediction
    Unc_y=np.zeros(len(x_g)) # Uncertainty in the prediction
    J_i_mean=np.zeros(n_e) # in sample error
    J_o_mean=np.zeros(n_e) # out of sample error
    w_e=np.zeros((3,n_e)) # Output distribution of weights
    # Take back the x vector
    x=X[:,1]
    # Loop over the ensemble
    for j in range(n_e):
        # Split the
        xs, xss, ys, yss = train_test_split(x,y, test_size=tp)
        # Fit the weights
        w_s=
        # Assign vectors to the distributions
        w_e[:,j]=w_s
        # Make in-sample prediction-----
        y_p_s=
        # In-sample error
        J_i_mean[j]=
        # Make out-of sample prediction (and errors)
        y_p_ss=
        # Out of sample error
        J_o_mean[j]=
        # Fill the population matrix
        Pop_Y[:,j]=

    # Compute the mean and the uncertainty
    mu_y=
    Var_y=
    Var_Y_model=
    Var_y_tot=
    Unc_y=1.96*np.sqrt(Var_y_tot)

    return mu_y, Unc_y, J_i_mean, J_o_mean, w_e

```

Here's a possible structure

The general idea is:

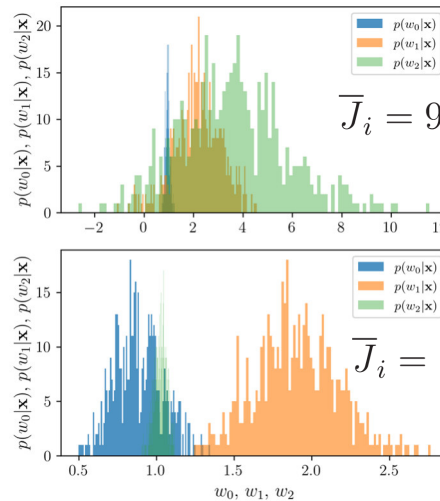
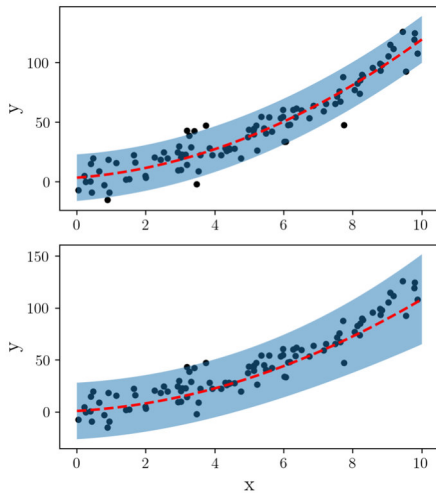
- (1) Initialize all the vectors that will be exported and the matrices that will be filled
- (2) Loop over the ensemble. Use scikit learn's function `train_test_split`
- (3) Train the regression on the training data and compute performances on both training and test data
- (4) Compute the ensemble statistics using python's function `np.mean` and `np.std`

Slide 7

Step 3: Uncertainty Analysis (2)



Here are the uncertainties and the distribution of weights for the ordinary and the regularized problem (taking $k_{\text{A max}} = 1000$)



Ordinary

$$\bar{J}_i = 95.3 \quad \bar{J}_o = 106.0$$

Regularized

$$\bar{J}_i = 196.4 \quad \bar{J}_o = 187.2$$

Slide 8

Step 4: Batch Gradient Descent



Write a python function to implement the mini-batch gradient descent, taking as inputs a function to compute costs and gradients, a decaying learning rate, a given number of epochs and a batch size.

First, you should write the following functions

```
def cost(w,X,y):
```

```
    return J
```

```
def grad(w,X,y):
```

```
    return Nabla_J
```

Then you should aim for a function of this form:

```
def Batch_GD(cost,grad,w0,X,y,eta_0,decay,n_epochs,n_batch):
```

```
    return w_opt, w_evolution, Err_SGD
```

Slide 9

```
def Batch_GD(cost,grad,w0,X,y,eta_0,decay,n_epochs,n_batch):
    # Prepare the loop per epoch
    n_iter=
    # number of points and features
    n_p,n_f=np.shape(X)
    # Initialize batch sample Design Matrix
    X_b=
    # Current estimate of w
    w=w0
    # Initialize the weight evolution and the error evolution
    Err_SGD=np.zeros(n_iter); #Err_SGD[0]=cost(w,X,y)
    w_evolution=np.zeros((n_f,n_iter)); #w_evolution[:,0]=w0

    for j in range(n_iter):
        # Select randomly some data points for the batch
        # Note that replace=False means that there is no repetition
        Indices=
        # Construct the matrix X_b
        X_b=X[Indices,:]; y_b=y[Indices]
        # Get the current cost
        Err_SGD[j]=cost(w,X_b,y_b)
        #Get the gradient
        Nabla_J_w=grad(w,X_b,y_b)
        # Compute the learning rate
        eta=eta_0/(1+decay*j)
        # Weight update
        w=
        # Store the result in the history
        w_evolution[:,j]=w
        # Message
        Mex='Iteration: {:d}; Epoch: {:d}; Cost: {:.3f}; Grad_abs: {:.3f}'\
            .format(j,j//n_batch,Err_SGD[j],np.linalg.norm(Nabla_J_w))
        print(Mex)

    w_opt=w # Final result on the weight
    return w_opt, w_evolution, Err_SGD
```



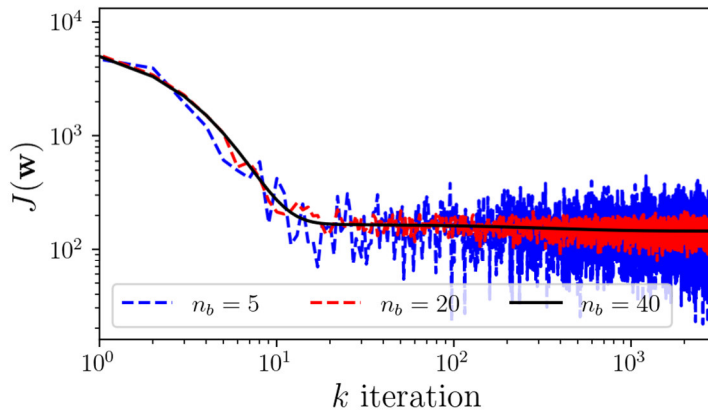
Here's a possible structure

The general idea is:

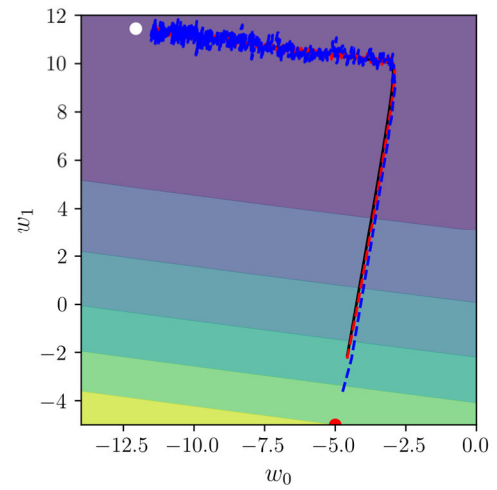
- (1) Take some random indices for the batch selection
- (2) Sample the data
- (3) Evaluate cost and gradients
- (4) Update the current estimate of the weights

Slide 10

Evolution of the Learning Process



Example results for the regression of a line using ordinary (unregularized) least square



Slide 11

The end!

Slide 12