

Hands on Machine Learning for Fluid Dynamics

7 – 11 February 2022




Lecture 5 **A Review of Optimization Tools**

Pedro Afonso Marques
pedro.marques@vki.ac.be


Introduction

Optimization


“The process of finding the best possible solution to a problem. In mathematics, this often consists of maximizing or minimizing the value of a certain function, perhaps subject to given constraints.” - Oxford dictionary




Curve fitting





System identification [1]



Stress field on NURBS elements (Iter.=1)



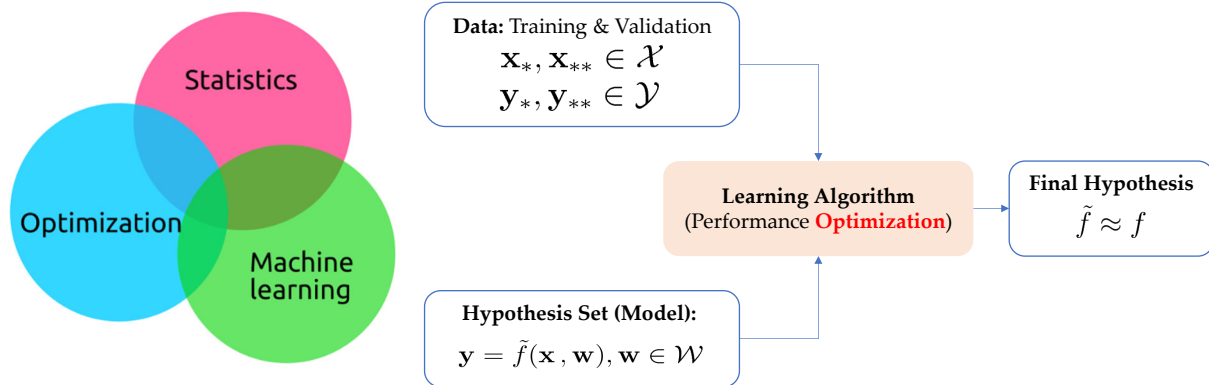
Shape optimization [2]



Economic models & resource allocation [3]

Introduction

Most Machine Learning algorithms can be framed as optimization problems...



Note: * denotes training data, ** denotes validation data

Table of Contents



1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
2. Overview of optimization methods
3. Gradient-based methods
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
4. Exercise: Mass-spring-damper system identification

Table of Contents



1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
2. Overview of optimization methods
3. Gradient-based methods
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
4. Exercise: Mass-spring-damper system identification

Formulating an optimization problem

Minimize:

- Objective function $f(\mathbf{x})$

Where:


- Vector of variables $\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_p \end{bmatrix}$

Subject to:

- Inequality constraints: $c_i(\mathbf{x}) \leq 0$ for $i = 1, \dots, m$
- Equality constraints: $c_e(\mathbf{x}) = 0$ for $e = 1, \dots, n$
- Variable range: $x_j \in [x_j^l, x_j^u]$ for $j = 1, \dots, p$

Example: Rosenbrock function

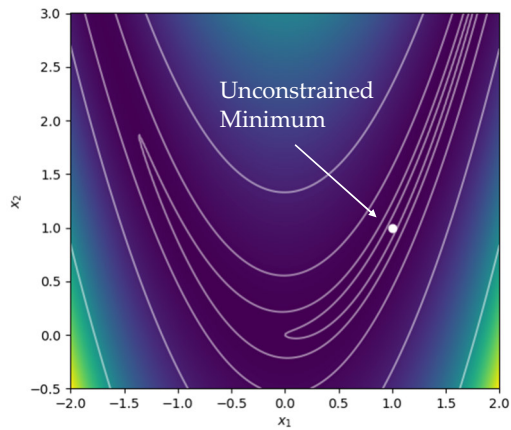
$$f(x_1, x_2) = (1 - x_1)^2 + 100(x_2 - x_1^2)^2$$




Formulating an optimization problem

Example: Rosenbrock function

$$f(x_1, x_2) = (1 - x_1)^2 + 100(x_2 - x_1^2)^2$$



Unconstrained minimum: $f(1,1) = 0$



Constrained minimum: $f(-0.9, 0.8) \approx 3.68$

Table of Contents



1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
2. Overview of optimization methods
3. Gradient-based methods
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
4. Exercise: Mass-spring-damper system identification

Constrained single-objective optimization

1. Handle constraints directly
2. External penalty method
3. Internal penalty method

Constrained single-objective optimization

1. Handle constraints directly


- Equality: Lagrange multiplier method
- Inequality: Linear/nonlinear/quadratic programming

2. External penalty method

$$\tilde{f}(\mathbf{x}) = f(\mathbf{x}) + R \cdot \sum_{i=1}^m \delta_i \cdot (c_i(\mathbf{x}))^2 \quad \text{where} \quad \delta_i = \begin{cases} 0 & , \quad c_i(\mathbf{x}) \leq 0 \\ 1 & , \quad c_i(\mathbf{x}) > 0 \end{cases}$$

3. Internal penalty method

$$\tilde{f}(\mathbf{x}) = f(\mathbf{x}) - R \cdot \sum_{i=1}^m \frac{1}{c_i(\mathbf{x})}$$



Constrained single-objective optimization

1. Handle constraints directly



- Equality: Lagrange multiplier method
- Inequality: Linear/nonlinear/quadratic programming

2. External penalty method

$$\tilde{f}(\mathbf{x}) = f(\mathbf{x}) + R \cdot \sum_{i=1}^m \delta_i \cdot (c_i(\mathbf{x}))^2 \quad \text{where} \quad \delta_i = \begin{cases} 0 & , \quad c_i(\mathbf{x}) \leq 0 \\ 1 & , \quad c_i(\mathbf{x}) > 0 \end{cases}$$

3. Internal penalty method

$$\tilde{f}(\mathbf{x}) = f(\mathbf{x}) - R \cdot \sum_{i=1}^m \frac{1}{c_i(\mathbf{x})}$$



Constrained single-objective optimization

1. Handle constraints directly

- Equality: Lagrange multiplier method
- Inequality: Linear/nonlinear/quadratic programming

2. External penalty method

$$\tilde{f}(\mathbf{x}) = f(\mathbf{x}) + R \cdot \sum_{i=1}^m \delta_i \cdot (c_i(\mathbf{x}))^2 \quad \text{where} \quad \delta_i = \begin{cases} 0 & , \quad c_i(\mathbf{x}) \leq 0 \\ 1 & , \quad c_i(\mathbf{x}) > 0 \end{cases}$$

3. Internal penalty method

$$\tilde{f}(\mathbf{x}) = f(\mathbf{x}) - R \cdot \sum_{i=1}^m \frac{1}{c_i(\mathbf{x})}$$




Table of Contents



- 1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
- 2. Overview of optimization methods
- 3. Gradient-based methods
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
- 4. Exercise: Mass-spring-damper system identification


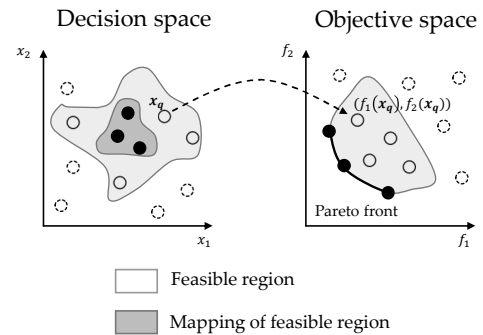
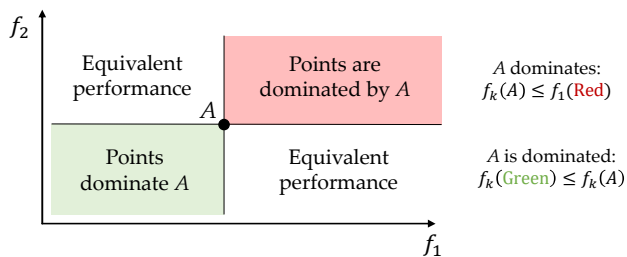
Multi-objective optimization

Minimize:

- Multiple objective functions $f_k(\mathbf{x})$ for $k = 1, \dots, l$

Subject to:

- Inequality constraints: $c_i(\mathbf{x}) \leq 0$ for $i = 1, \dots, m$
- Equality constraints: $c_e(\mathbf{x}) = 0$ for $e = 1, \dots, n$
- Variable range: $x_j \in [x_j^l, x_j^u]$ for $j = 1, \dots, p$



Difference between objectives and constraints

Single-objective constrained optimization

E.g.: Optimize the volume $V(x, y, z)$ of a component, but make sure that the maximum stress $\sigma(x, y, z)$ is lower than a certain amount.

$$\begin{aligned} \min \quad & V(\mathbf{x}) \\ \text{s.t.} \quad & \sigma(\mathbf{x}) \leq C \\ & \mathbf{x} \in \mathbf{X} \end{aligned}$$

Multi-objective optimization

E.g.: Simultaneously optimize the volume $V(x, y, z)$ of a component, and minimize maximum stress $\sigma(x, y, z)$ it is subjected to.

$$\begin{aligned} \min \quad & (V(\mathbf{x}), \sigma(\mathbf{x})) \\ \text{s.t.} \quad & \mathbf{x} \in \mathbf{X} \end{aligned}$$

Multi to single-objective optimization

1. Weighted sum of each individual objective
2. Min-max formulation
3. Constrained minimum

Multi to single-objective optimization

1. Weighted sum of each individual objective


$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$



Multi to single-objective optimization

1. Weighted sum of each individual objective


$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$



Multi to single-objective optimization

1. Weighted sum of each individual objective


$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$



Multi to single-objective optimization

1. Weighted sum of each individual objective


$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$



Multi to single-objective optimization

1. Weighted sum of each individual objective

$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$

Multi to single-objective optimization

1. Weighted sum of each individual objective


$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$



Multi to single-objective optimization

1. Weighted sum of each individual objective


$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$



Multi to single-objective optimization

1. Weighted sum of each individual objective


$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$



Multi to single-objective optimization

1. Weighted sum of each individual objective

$$F(\mathbf{x}) = \sum_{k=1}^l w_k f_k(\mathbf{x})$$

2. Min-max formulation

$$F(\mathbf{x}) = \max(\omega_1 f_1(\mathbf{x}), \omega_2 f_2(\mathbf{x}), \dots, \omega_l f_l(\mathbf{x}))$$

3. Constrained minimum

- Minimize f_1 assuming $f_2 = \text{cte}$
- Minimize f_2 assuming $f_1 = \text{cte}$


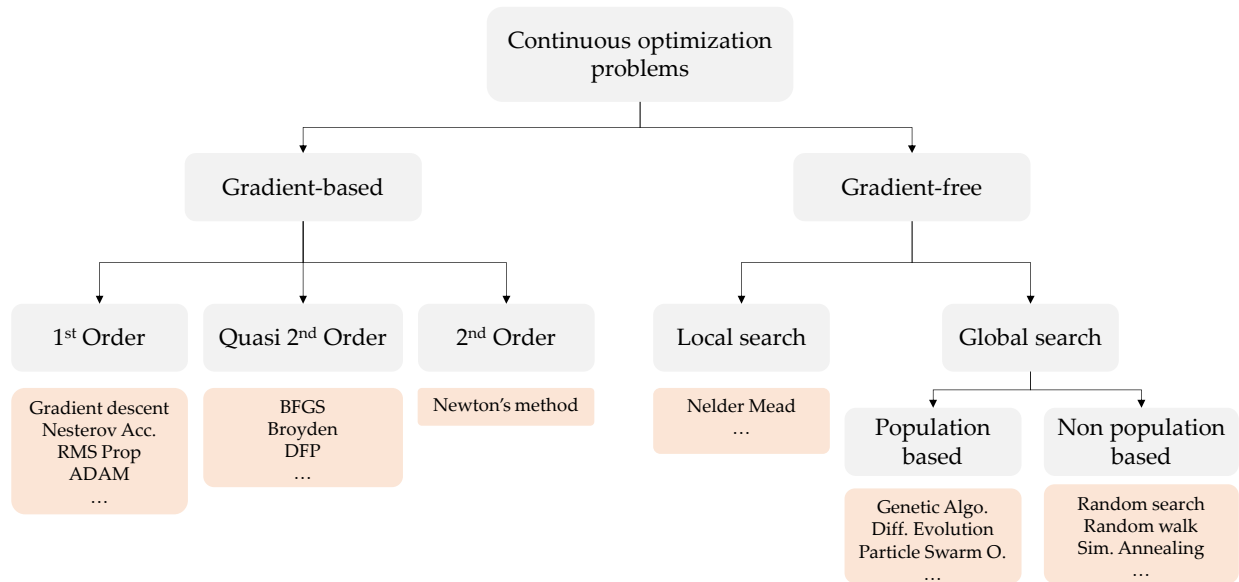


Table of Contents

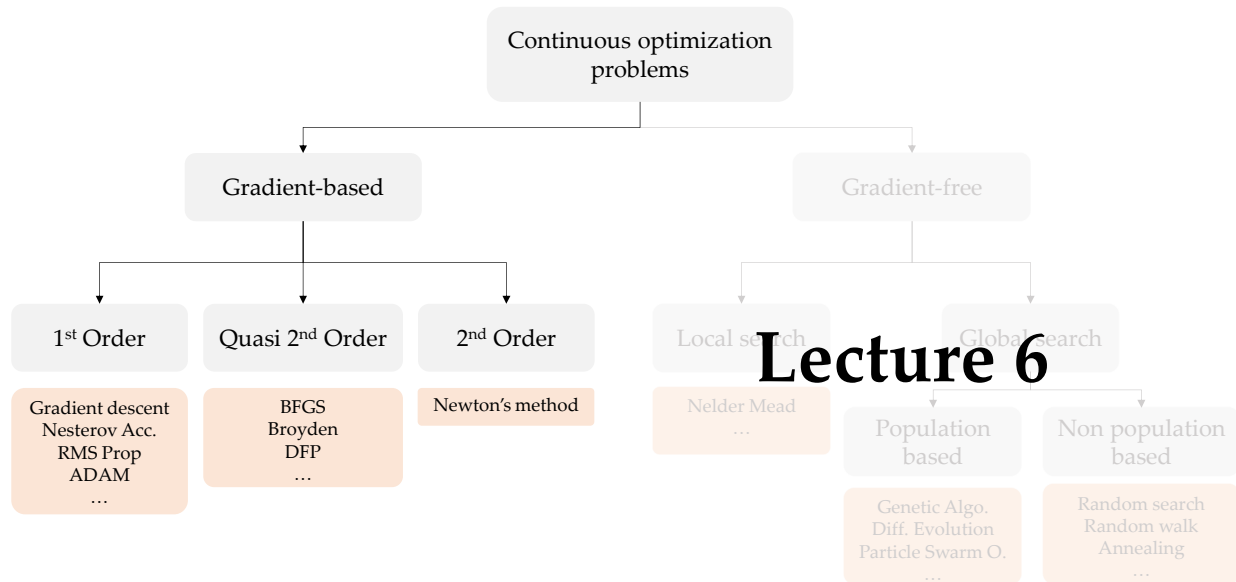


- 1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
- 2. Overview of optimization methods
- 3. Gradient-based methods
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
- 4. Exercise: Mass-spring-damper system identification

Optimization methods



Optimization methods



Lecture 6

Table of Contents



- 1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
- 2. Overview of optimization methods
- 3. Gradient-based methods**
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
- 4. Exercise: Mass-spring-damper system identification

Gradient-based methods

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta^{(k)} d^{(k)}$$

First order steepest descent

- Search direction is the steepest descent
- Cheap computational evaluation, only requires the gradient

Conjugate gradient methods

- Search direction is **not** the steepest descent
- Perform line search for optimal learning rate

High order gradient descent

- Compute or estimate the Hessian matrix
- Converge in fewer iterations at cost of computational load

Table of Contents



- 1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
- 2. Overview of optimization methods
- 3. Gradient-based methods**
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
- 4. Exercise: Mass-spring-damper system identification

Gradient descent methods

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \eta(k) \nabla_w J(\mathbf{w}; \mathbf{X}, \mathbf{y})$$

ANN Optimizers

Modify the learning rate:

$$\eta^{(k)}$$

'Classic' gradient descent methods

- Batch gradient descent
- Mini-batch gradient descent
- ...

Table of Contents



- 1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
- 2. Overview of optimization methods
- 3. Gradient-based methods**
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients**
 - 3.3. Quasi-Newton methods
- 4. Exercise: Mass-spring-damper system identification

Nonlinear Conjugate Gradient method

1. Define search direction

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta^{(k)} d^{(k)}$$

- In linear problems, we search in orthogonal directions.
- It can be shown that these models converge in \mathbf{n} iterations, where \mathbf{n} is the dimension of the search space!

$$d^{(k+1)} = \underbrace{-\nabla_w J(\mathbf{w}^{(k+1)})}_{\text{Steepest descent direction}} + \underbrace{\beta^{(k+1)} d^{(k)}}_{\text{Previous search direction}}$$

Fletcher-Reeves:

$$\beta^{(k+1)} = \frac{\nabla_w J(\mathbf{w}^{(k+1)})^T \nabla_w J(\mathbf{w}^{(k+1)})}{\nabla_w J(\mathbf{w}^{(k)})^T \nabla_w J(\mathbf{w}^{(k)})}$$

Polak-Ribière:

$$\beta_{k+1} = \frac{\nabla_w J(\mathbf{w}^{(k+1)})^T (\nabla_w J(\mathbf{w}^{(k+1)}) - \nabla_w J(\mathbf{w}^{(k)}))}{\|\nabla_w J(\mathbf{w}^{(k)})\|^2}$$

Nonlinear Conjugate Gradient method

2. Find optimal step length

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta^{(k)} d^{(k)}$$

- Minimize along search direction $\min_{\eta} J(\mathbf{w}^{(k)} + \eta d^{(k)})$
- Iterative algorithm for 1D minimizer (e.g. cubic interpolation)
- Satisfy “ideal” step length conditions (e.g. Wolfe, Goldstein)

Sufficient decrease:

$$J(\mathbf{w}^{(k)} + \eta d^{(k)}) \leq J(\mathbf{w}^{(k)}) + c_1 \eta \nabla_w J(\mathbf{w}^{(k)})^T d^{(k)}$$

Curvature condition:

$$\left| \nabla_w J(\mathbf{w}^{(k)} + \eta d^{(k)})^T d_k \right| \leq c_2 \left| \nabla_w J(\mathbf{w}^{(k)})^T d_k \right|$$

The exact minimizer \mathbf{w}^* can be found in one iteration if f is a quadratic function

Rewrite the objective function as:

$$J(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T Q \mathbf{w} - b^T \mathbf{w}$$

Using the steepest descent algorithm:

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta^{(k)} \nabla_w J(\mathbf{w}^{(k)})$$

The exact minimizer is found:

$$\nabla_{\eta} J(\mathbf{w}^{(k)} - \eta^* \nabla_w J(\mathbf{w}^{(k)})) = 0$$

$$\eta^* = \frac{\nabla_w J(\mathbf{w}^{(k)})^T \nabla_w J(\mathbf{w}^{(k)})}{\nabla_w J(\mathbf{w}^{(k)})^T Q \nabla_w J(\mathbf{w}^{(k)})}$$

Table of Contents



- 1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
- 2. Overview of optimization methods
- 3. Gradient-based methods**
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods**
- 4. Exercise: Mass-spring-damper system identification

Quasi-Newton methods

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \eta^{(k)} \underbrace{\left(B^{(k)} \right)^{-1}}_{\approx \text{Hessian}} \nabla_w J(\mathbf{w})$$

Quadratic approximation of objective function:

$$m^{(k)}(d) = J(\mathbf{w}^{(k)}) + \nabla_w J(\mathbf{w}^{(k)})^T \underline{d} + \frac{1}{2} \underline{d}^T B^{(k)} \underline{d}$$

Search direction
 $d = -B^{-1} \nabla_w J(\mathbf{w})$

- Gradient of the quadratic model must match $\nabla_w J(\mathbf{w})$ for $w^{(k)}$ and $w^{(k+1)}$:

$$B^{(k+1)} \left(\mathbf{w}^{(k+1)} - \mathbf{w}^{(k)} \right) = \nabla J_w \left(\mathbf{w}^{(k+1)} \right) - \nabla J_w \left(\mathbf{w}^{(k)} \right)$$

- Can be solved if:

$$\left(\mathbf{w}^{(k+1)} - \mathbf{w}^{(k)} \right)^T \nabla J_w \left(\mathbf{w}^{(k+1)} \right) - \nabla J_w \left(\mathbf{w}^{(k)} \right) > 0 \longrightarrow \text{Line search for } \eta^{(k)} \text{ with Strong Wolfe conditions}$$

Quasi-Newton methods

1. Compute search direction: $d^{(k)} = -H^{(k)} \nabla_w J(\mathbf{w})$ $s^{(k)} = x^{(k+1)} - x^{(k)}$
2. Line search with Strong Wolfe conditions: $\eta^{(k)}$ $y^{(k)} = \nabla_w J(\mathbf{w}^{(k+1)}) - \nabla_w J(\mathbf{w}^{(k)})$
3. Compute next solution: $\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta^{(k)} d^{(k)}$ $\rho^{(k)} = \left(y^{(k)} s^{(k)} \right)^{-1}$
4. Compute $H^{(k+1)}$:

$$\text{(DFP)} \quad H^{(k+1)} = H^{(k)} - \frac{H^{(k)} y^{(k)} (y^{(k)})^T H^{(k)}}{(y^{(k)})^T H^{(k)} y^{(k)}} + \frac{s^{(k)} (s^{(k)})^T}{(y^{(k)})^T s^{(k)}}$$

$$\text{(BFGS)} \quad H^{(k+1)} = \left(I - \rho^{(k)} s^{(k)} (y^{(k)})^T \right) H^{(k)} \left(I - \rho^{(k)} y^{(k)} (s^{(k)})^T \right) + \rho^{(k)} s^{(k)} (s^{(k)})^T$$

Initial estimate for the Hessian?

Finite difference estimate

Identity matrix

...

Table of Contents



- 1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
- 2. Overview of optimization methods
- 3. Gradient-based methods
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
- 4. Exercise: Mass-spring-damper system identification

Gradient-based optimization: Minimize Rosenbrock function

- Content here...

Table of Contents




1. Formulating an optimization problem
 - 1.1. Constrained optimization
 - 1.2. Multi-objective optimization
2. Overview of optimization methods
3. Gradient-based methods
 - 3.1. Gradient descent
 - 3.2. Conjugate gradients
 - 3.3. Quasi-Newton methods
4. Exercise: Mass-spring-damper system identification

Exercise: Mass-spring-damper system identification

Steepest descent
BFGS (compare with scipy)
Newton
Nelder Mead

Compare:

- Total elapsed time
- Cost-function vs iterations
- Gradient = efficiency



2nd order system
Step response

$$m\ddot{x} + b\dot{x} + kx = 0$$

Title

Content here...

Gradient-based methods

Machine learning

Heuristic learning rate.

Keep the learning rate,

Add momentum, rescaling,

(connection with Jan in lecture 11)

See **DDFM Lecture 10 Slide 29**

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta^{(k)} d^{(k)}$$

ANN Optimizers

Modify the learning rate:

$$\eta^{(k)}$$

Conjugate gradient methods

- Search direction is **not** the steepest descent
- Perform line search for optimal learning rate

High order gradient descent

- Compute the Hessian (or approximate)
- Converge in fewer iterations...

State that the gradient is available through a function (for these cases assume that it's possible)

To do

Func(cost_function,grad_cost_function,n_iter)

Exercise:

- Gradient descent
- Quasi-Newton (BFGS)
- Newton
- Compare convergence rate (plot residual vs iterations)
- Compare time called by our routine vs scipy
- What if we don't give the function for the gradient
 - Takes even longer, because it has to evaluate it

```
Miguel Alfonso Mendez 16:17
x_star,conv_h=Opt(cost_fun,grad_func,n_iter)

x_star,conv_h=Opt(cost_fun,grad_fun,100)

Opt (1) plain gradient descent
Opt (2) gradient descent with line search
Opt (3) BFGS
```

Plot convergence rate
Compare evaluation time
Our evaluation vs scipy

Give a template with some missing parts, which can then be filled together with the people

Nonlinear Conjugate Gradient method

1. Define search direction

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta^{(k)} \vec{d}^{(k)}$$

2. Find optimal step length

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta^{(k)} \vec{d}^{(k)}$$

- Minimize in the search direction $\min_{\mathbf{w}} J(\mathbf{w}^{(k)} + \eta^{(k)} \vec{d}^{(k)})$
- Iterative algorithm for 1D minimizer (e.g. cubic interpolation)
- Satisfy “ideal” step length conditions (e.g. Wolfe, Goldstein)

Sufficient decrease: $f(x_k + \alpha \vec{p}_k) \leq f(x_k) + c_1 \alpha \nabla f_k^T p_k$ Curvature condition: $|\nabla f(x_k + \alpha \vec{p}_k)^T p_k| \leq c_2 |\nabla f_k^T p_k|$



scipy.optimize.line_search

scipy.optimize.line_search(f, myfprime, x0, p0, gfk=None, old_fval=None, old_old_fval=None, args=(), c1=0.0001, c2=0.9, amax=None, extra_condition=None, maxiter=10) [source]
Find alpha that satisfies strong Wolfe conditions.

The exact minimizer \mathbf{w}^* can be found in one iteration if f is a quadratic function

Rewrite the objective function as:

$$f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T Q \mathbf{x} - b^T \mathbf{x}$$

Using the steepest descent algorithm:

$$x_{k+1} = x_k - \alpha_k \nabla f_k$$

The exact minimizer is found:

$$\nabla_{\alpha} f(x_k - \alpha^* \nabla f_k) = 0$$

$$\alpha^* = \frac{\nabla f_k^T \nabla f_k}{\nabla f_k^T Q \nabla f_k}$$

Linear conjugate gradient method

Minimize quadratic function: $\phi(\mathbf{x}) = \frac{1}{2}\mathbf{x}^T A \mathbf{x} - b^T \mathbf{x}$

Residual defined as the gradient: $\nabla \phi(\mathbf{x}) = A\mathbf{x} - b = r(\mathbf{x})$

Search directions p_k : conjugate vector wrt A $p_i^T A p_j = 0 \quad \forall \quad i \neq j$

Initial conditions: $r_0 = A\mathbf{x}_0 - b$ $p_0 = r_0$ Exact minimizer for quadratic functions

Minimization algorithm:

- Ideal step length through line search $\alpha_k = -\frac{r_k^T p_k}{p_k^T A p_k}$ Exact minimizer for quadratic functions
- Move to the minimizer in direction p_k $x_{k+1} = x_k + \alpha_k \nabla f_k$
- Calculate new residual r_{k+1} $r_{k+1} = A x_{k+1} - b$
- New direction p_{k+1} conjugate wrt A $p_{k+1} = -r_{k+1} + \beta_{k+1} p_k$
Linear combination of steepest descent $-r_{k+1}$ and previous search direction $\beta_{k+1} = \frac{r_{k+1}^T A p_k}{p_k^T A p_k}$ Constant enforces that $p_{k-1}^T A p_k = 0$, so that they are conjugate directions

How to choose the conjugate direction set $\{p_0, p_1, \dots, p_{n-1}\}$ for matrix A $[n \times n]$?

Eigenvectors v_1, v_2, \dots, v_n of A

- Mutually orthogonal (linearly independent)
 $\langle v_i, v_j \rangle = v_j^T v_i = 0 \quad \forall \quad i \neq j$
- Conjugate with respect to A
 $A v_i = \lambda v_i \Leftrightarrow v_j^T A v_i = \lambda v_j^T v_i \Rightarrow v_j^T A v_i = 0$
- Too expensive for large-scale applications

More efficient approaches are used (see left)

Linear conjugate gradient method

If A is a diagonal matrix $[n \times n]$, we converge to the minimizer of $\phi(x)$ in n iterations.

We need a diagonal matrix to guarantee convergence!


$$\hat{x} = S^{-1}x \quad \text{where} \quad S = [p_0 \ p_1 \ \dots \ p_{n-1}]$$

$$\hat{\phi}(\hat{x}) = \frac{1}{2} \hat{x}^T (S^T A S) \hat{x} - (S^T b)^T \hat{x}$$

Since $S^T A S$ is diagonal, we minimize $\hat{\phi}$ with n iterations in the \hat{x} space, and then compute the true minimum:

$$x = S\hat{x}$$

We can create a symmetric matrix during preconditioning.



Successive minimization for a general convex quadratic function using the linear conjugate gradient method

Optimization methods

Gradient-based (local)

- Line search methods
 - Deepest descent
 - Newton's method
 - Quasi-newton methods
- Trust region methods
 - Quasi-newton methods
 - Newton's method

Mention this as reference:

<https://docs.scipy.org/doc/scipy/tutorial/optimize.html>

Gradient-free (global)

- Random search (simplest approach of all)
- Random walk (kind of a line search but random)
- Nelder-Mead (polytopes)
- Evolutionary algorithms (discussed in Lecture 6/7)
- Particle Swarm Optimization (discussed in lecture 6)

Hybrid

- Bayesian optimization (following lecture)

This lecture: mainly gradient based. Apply them to some problems.

Exercise: Fitting a turbulent boundary layer

Non-Linear Least Squares (only scipy)

Steepest descent



BFGS (compare with scipy)

Newton

Nelder Mead

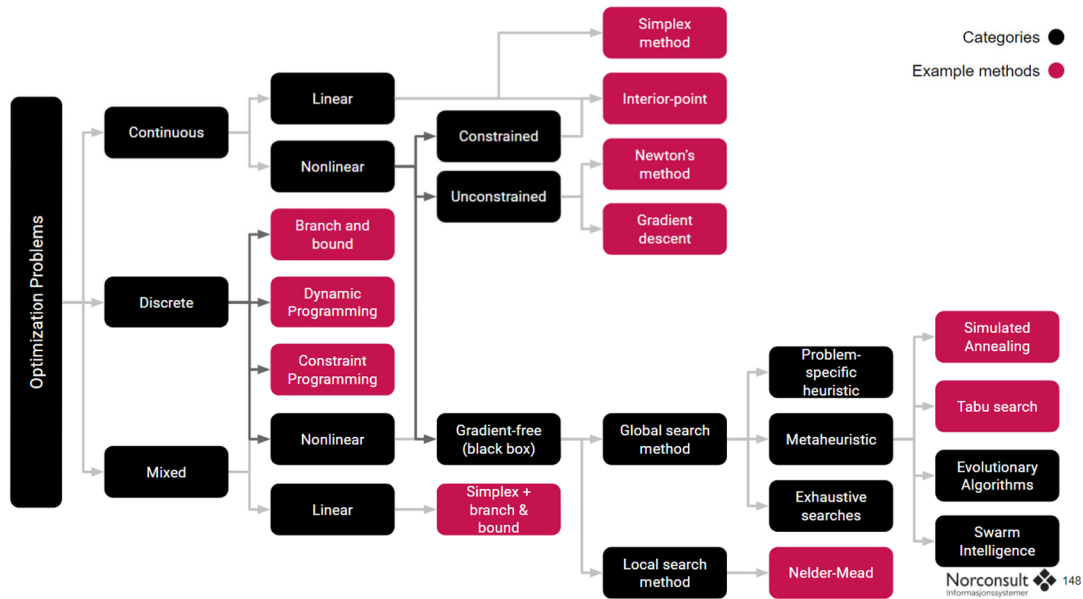
Compare:

- Total elapsed time
- Cost-function vs iterations
- Gradient = efficiency



$$u^+ = a_1 \log(1 + y^+ \kappa) + a_2 \left(1 - e^{-\frac{y^+}{11}} - \frac{y^+}{11} e^{-\frac{y^+}{3}} \right)$$

Optimization methods



Unused figures

Single-objective

